

# Lab 4 & 5

Daniel Mehta

Part 1a: Follow the tutorial code given in the lecture  
(<https://www.tensorflow.org/tutorials/genera>



```
In [5]: import tensorflow as tf
        tf.__version__
```

```
Out[5]: '2.19.0'
```

```
In [6]: import glob
        import imageio
        import matplotlib.pyplot as plt
        import numpy as np
        import os
        import PIL
        from tensorflow.keras import layers
        import time

        from IPython import display
```

```
In [7]: (train_images, train_labels), (_, _) = tf.keras.datasets.mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11490434/11490434 ————— 1s 0us/step

```
In [8]: train_images = train_images.reshape(train_images.shape[0], 28, 28, 1).astype('float')
        train_images = (train_images - 127.5) / 127.5 # Normalize the images to [-1, 1]
```

```
In [9]: BUFFER_SIZE = 60000
        BATCH_SIZE = 256
```

```
In [10]: # Batch and shuffle the data
         train_dataset = tf.data.Dataset.from_tensor_slices(train_images).shuffle(BUFFER_SIZE)
```

```
In [11]: def make_generator_model():
        model = tf.keras.Sequential()
        model.add(layers.Dense(7*7*256, use_bias=False, input_shape=(100,)))
        model.add(layers.BatchNormalization())
        model.add(layers.LeakyReLU())

        model.add(layers.Reshape((7, 7, 256)))
        assert model.output_shape == (None, 7, 7, 256) # Note: None is the batch size
```

```

model.add(layers.Conv2DTranspose(128, (5, 5), strides=(1, 1), padding='same', u
assert model.output_shape == (None, 7, 7, 128)
model.add(layers.BatchNormalization())
model.add(layers.LeakyReLU())

model.add(layers.Conv2DTranspose(64, (5, 5), strides=(2, 2), padding='same', us
assert model.output_shape == (None, 14, 14, 64)
model.add(layers.BatchNormalization())
model.add(layers.LeakyReLU())

model.add(layers.Conv2DTranspose(1, (5, 5), strides=(2, 2), padding='same', use
assert model.output_shape == (None, 28, 28, 1)

return model

```

```

In [12]: generator = make_generator_model()

noise = tf.random.normal([1, 100])
generated_image = generator(noise, training=False)


plt.imshow(generated_image[0, :, :, 0], cmap='gray')

```

C:\Users\danie\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\layers\core\dense.py:93: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
Out[12]: <matplotlib.image.AxesImage at 0x16d5f1bbcd0>
```



```
In [13]: def make_discriminator_model():
    model = tf.keras.Sequential()
    model.add(layers.Conv2D(64, (5, 5), strides=(2, 2), padding='same',
                             input_shape=[28, 28, 1]))

    model.add(layers.LeakyReLU())
    model.add(layers.Dropout(0.3))

    model.add(layers.Conv2D(128, (5, 5), strides=(2, 2), padding='same'))
    model.add(layers.LeakyReLU())
    model.add(layers.Dropout(0.3))

    model.add(layers.Flatten())
    model.add(layers.Dense(1))

    return model
```

```
In [14]: discriminator = make_discriminator_model()
    decision = discriminator(generated_image)
    print (decision)
```

```
tf.Tensor([[ -0.00176129]], shape=(1, 1), dtype=float32)
```

C:\Users\danie\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\layers\convolutional\base\_conv.py:113: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [15]: # This method returns a helper function to compute cross entropy loss
    cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=True)
```

```
In [16]: def discriminator_loss(real_output, fake_output):
    real_loss = cross_entropy(tf.ones_like(real_output), real_output)
    fake_loss = cross_entropy(tf.zeros_like(fake_output), fake_output)
    total_loss = real_loss + fake_loss
    return total_loss
```

```
In [17]: def generator_loss(fake_output):
    return cross_entropy(tf.ones_like(fake_output), fake_output)
```

```
In [18]: generator_optimizer = tf.keras.optimizers.Adam(1e-4)
    discriminator_optimizer = tf.keras.optimizers.Adam(1e-4)
```

```
In [19]: checkpoint_dir = './training_checkpoints'
    checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt")
    checkpoint = tf.train.Checkpoint(generator_optimizer=generator_optimizer,
                                     discriminator_optimizer=discriminator_optimizer,
                                     generator=generator,
                                     discriminator=discriminator)
```

```
In [26]: #EPOCHS = 50
    EPOCHS = 10 # Lowered to ten for time
    noise_dim = 100
    num_examples_to_generate = 16
```

```
# You will reuse this seed overtime (so it's easier)
# to visualize progress in the animated GIF)
seed = tf.random.normal([num_examples_to_generate, noise_dim])
```

```
In [21]: # Notice the use of `tf.function`
# This annotation causes the function to be "compiled".
@tf.function
def train_step(images):
    noise = tf.random.normal([BATCH_SIZE, noise_dim])

    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)

        real_output = discriminator(images, training=True)
        fake_output = discriminator(generated_images, training=True)

        gen_loss = generator_loss(fake_output)
        disc_loss = discriminator_loss(real_output, fake_output)

        gradients_of_generator = gen_tape.gradient(gen_loss, generator.trainable_variables)
        gradients_of_discriminator = disc_tape.gradient(disc_loss, discriminator.trainable_variables)

        generator_optimizer.apply_gradients(zip(gradients_of_generator, generator.trainable_variables))
        discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator, discriminator.trainable_variables))
```

```
In [22]: def train(dataset, epochs):
    for epoch in range(epochs):
        start = time.time()

        for image_batch in dataset:
            train_step(image_batch)

        # Produce images for the GIF as you go
        display.clear_output(wait=True)
        generate_and_save_images(generator,
                                epoch + 1,
                                seed)

        # Save the model every 15 epochs
        if (epoch + 1) % 15 == 0:
            checkpoint.save(file_prefix = checkpoint_prefix)

        print('Time for epoch {} is {} sec'.format(epoch + 1, time.time()-start))

    # Generate after the final epoch
    display.clear_output(wait=True)
    generate_and_save_images(generator,
                            epochs,
                            seed)
```

```
In [23]: def generate_and_save_images(model, epoch, test_input):
    # Notice `training` is set to False.
    # This is so all layers run in inference mode (batchnorm).
    predictions = model(test_input, training=False)
```

```
fig = plt.figure(figsize=(4, 4))

for i in range(predictions.shape[0]):
    plt.subplot(4, 4, i+1)
    plt.imshow(predictions[i, :, :, 0] * 127.5 + 127.5, cmap='gray')
    plt.axis('off')

plt.savefig('image_at_epoch_{:04d}.png'.format(epoch))
plt.show()
```

In [24]: `train(train_dataset, EPOCHS)`




In [25]: `checkpoint.restore(tf.train.latest_checkpoint(checkpoint_dir))`

Out[25]: `<tensorflow.python.checkpoint.checkpoint.CheckpointLoadStatus at 0x16d605772e0>`

In [27]: `# Display a single image using the epoch number`  
`def display_image(epoch_no):`  
 `return PIL.Image.open('image_at_epoch_{:04d}.png'.format(epoch_no))`

In [29]: `display_image(50)`

Out[29]:



```
In [30]: anim_file = 'drgan.gif'

with imageio.get_writer(anim_file, mode='I') as writer:
    filenames = glob.glob('image*.png')
    filenames = sorted(filenames)
    for filename in filenames:
        image = imageio.imread(filename)
        writer.append_data(image)
    image = imageio.imread(filename)
    writer.append_data(image)
```

C:\Users\danie\AppData\Local\Temp\ipykernel\_18276\1982054950.py:7: DeprecationWarning: Starting with ImageIO v3 the behavior of this function will switch to that of `imageio.v3.imread`. To keep the current behavior (and make this warning disappear) use `import imageio.v2 as imageio` or call `imageio.v2.imread` directly.


```
    image = imageio.imread(filename)
```

C:\Users\danie\AppData\Local\Temp\ipykernel\_18276\1982054950.py:9: DeprecationWarning: Starting with ImageIO v3 the behavior of this function will switch to that of `imageio.v3.imread`. To keep the current behavior (and make this warning disappear) use `import imageio.v2 as imageio` or call `imageio.v2.imread` directly.

```
    image = imageio.imread(filename)
```

```
In [31]: import tensorflow_docs.vis.embed as embed
         embed.embed_file(anim_file)
```

Out[31]:



## Part 1b: choose areal-life problem where GANs can be used as a full or part of the solution

### Dataset: Fashion MNIST

Retailers often need to generate synthetic fashion items for prototyping or virtual fitting rooms

```
In [34]: (train_images, train_labels), (_, _) = tf.keras.datasets.fashion_mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
29515/29515 ————— 0s 1us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
26421880/26421880 ————— 2s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
5148/5148 ————— 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
4422102/4422102 ————— 0s 0us/step
```

```
In [35]: train_images = train_images.reshape(train_images.shape[0], 28, 28, 1).astype('float')
train_images = (train_images - 127.5) / 127.5
```

```
BUFFER_SIZE = 60000  
BATCH_SIZE = 256
```


```
In [36]: train_dataset = tf.data.Dataset.from_tensor_slices(train_images).shuffle(BUFFER_SIZE)
```

```
In [37]: make_generator_model()  
         make_discriminator_model()
```

```
Out[37]: <Sequential name=sequential_3, built=True>
```

```
In [39]: EPOCHS = 10  
         noise_dim = 100  
         num_examples_to_generate = 16  
         seed = tf.random.normal([num_examples_to_generate, noise_dim])
```


```
In [40]: train(train_dataset, EPOCHS)
```



```
In [41]: display_image(EPOCHS)
```



Out[41]:



## Part 2: Evaluating Generator Performance

```
In [42]: gen_losses = []
         disc_losses = []
```

```
In [43]: @tf.function
         def train_step(images):
             noise = tf.random.normal([BATCH_SIZE,100])

             with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
                 generated_images = generator(noise, training=True)
                 real_output = discriminator(images, training=True)
                 fake_output = discriminator(generated_images, training=True)

                 gen_loss = generator_loss(fake_output)
                 disc_loss = discriminator_loss(real_output, fake_output)

                 gradients_of_generator = gen_tape.gradient(gen_loss, generator.trainable_variables)
                 gradients_of_discriminator = disc_tape.gradient(disc_loss, discriminator.trainable_variables)

                 generator_optimizer.apply_gradients(zip(gradients_of_generator, generator.trainable_variables))
                 discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator, discriminator.trainable_variables))

             return gen_loss, disc_loss
```

```
In [48]: def train(dataset, epochs):
         for epoch in range(epochs):
             start = time.time()
             gen_total = 0
```

```

disc_total = 0
batches = 0

for image_batch in dataset:
    gen_loss, disc_loss = train_step(image_batch)
    gen_total += gen_loss
    disc_total += disc_loss
    batches += 1

gen_losses.append(gen_total / batches)
disc_losses.append(disc_total / batches)

display.clear_output(wait=True)
generate_and_save_images(generator, epoch + 1, seed)

print(f'Epoch {epoch + 1}, Gen Loss: {gen_losses[-1]:.4f}, Disc Loss: {disc_
print(f'Time for epoch {epoch + 1} is {time.time() - start:.2f} sec')

display.clear_output(wait=True)
generate_and_save_images(generator, epochs, seed)


```

```

In [49]: EPOCHS = 10
seed = tf.random.normal([16, 100])

train(train_dataset, EPOCHS)


```



```

In [50]: plt.plot(gen_losses, label='Generator Loss')
plt.plot(disc_losses, label='Discriminator Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training Loss Curves')
plt.legend()
plt.show()

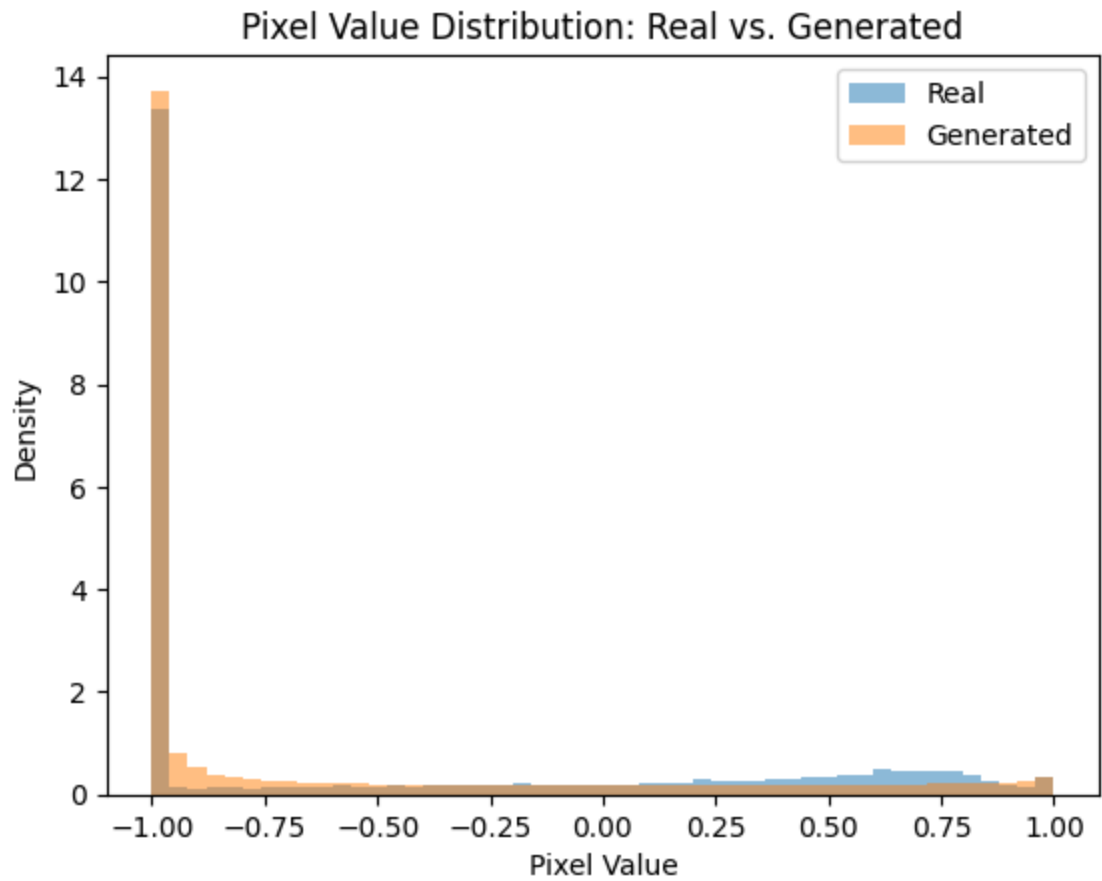
```



```
In [51]: #sample 1000 generated images
generated = generator(tf.random.normal([1000, 100]), training=False).numpy()
generated = generated.reshape(-1)

#Sample 1000 real images
real = train_images[:1000].reshape(-1)

#histogram
plt.hist(real, bins=50, alpha=0.5, label='Real', density=True)
plt.hist(generated, bins=50, alpha=0.5, label='Generated', density=True)
plt.title('Pixel Value Distribution: Real vs. Generated')
plt.xlabel('Pixel Value')
plt.ylabel('Density')
plt.legend()
plt.show()
```



In [ ]: