

Graduate Certificate in Artificial Intelligence with Machine Learning  
AIGC 5504 – Emerging Technologies in Artificial Intelligence

## Lab 10: Implementing Q-Learning for Reinforcement Learning

### Submission guidelines:

- For this lab, you will need to submit 1 PDF file.
  - Convert your codes to PDF.
  - Name the PDF as follows: firstname\_lastname\_LAB 10.pdf
  - Go to the course Blackboard → Labs folder → Lab Exercises 10 and submit the pdf.
- 

### Lab goals:

- Implement Q-Learning to solve a basic reinforcement learning problem.
  - Understand key RL concepts like states, actions, rewards, and policies.
  - Understand how neural networks approximate Q-values.
  - Explore the impact of different parameters (e.g., learning rate, discount factor).
- 

#### Part 1: Understanding the Problem

1. **Scenario:**
  - A robot is navigating a 5x5 grid to reach a goal while avoiding obstacles.
  - Each grid cell is a state, and actions are moving up, down, left, or right.
2. **Objective:**
  - Train a neural network to approximate the Q-values for state-action pairs and find the optimal path to the goal.

#### Part 2: Setting Up the Environment

1. **Grid Environment:**
  - Create a 5x5 grid where:
    - Top-left corner (0, 0) is the start.
    - Bottom-right corner (4, 4) is the goal.
    - Randomly place 3 obstacles.
2. **Rewards:**
  - +10 for reaching the goal.
  - -10 for hitting an obstacle.
  - 0 for all other states.
3. **State Representation:**
  - Represent states as one-hot encoded vectors or scaled grid coordinates (e.g., normalized (x, y) coordinates).

**Part 3: Implementing Deep Q-Learning****1. Neural Network Architecture:**

- Input: State representation (5x5 grid flattened to 25 inputs or (x, y) coordinates).
- Hidden Layers: 2 fully connected layers with ReLU activation (e.g., 64 neurons each).
- Output: Q-values for each action (4 outputs for up, down, left, right).

**2. Algorithm Steps:**

- Initialize the neural network for Q-value approximation.
- Use experience replay to store and sample experiences  $(s, a, r, s')$ .
- Implement  $\epsilon$ -greedy policy for exploration.
- Train the network using the Bellman equation
- Update the network weights to minimize the loss

**3. Steps:**

- Initialize replay memory (e.g., capacity of 10,000 experiences).
- For each episode:
  1. Start at the initial state.
  2. Select an action using  $\epsilon$ -greedy policy.
  3. Take the action and observe the next state and reward.
  4. Store the transition  $(s, a, r, s')$  in replay memory.
  5. Sample a batch of transitions from memory.
  6. Train the network to minimize loss.
  7. Update the target network periodically (e.g., every 10 episodes).

**Part 4: Visualizing Results****1. Performance Metrics:**

- Plot total reward per episode.
- Visualize the optimal path from start to goal on the grid.

**2. Q-Value Approximation:**

- Show the predicted Q-values for some states.

**Part 5: Reflection Questions**

Answer the following in your PDF submission:

1. How does using neural networks improve over traditional Q-tables?
2. What challenges did you face when implementing the Deep Q-Learning algorithm?
3. How do hyperparameters affect training?
4. Suggest a real-world application of Deep Reinforcement Learning and explain its implementation.

Enjoy!

---